

# Training Naive Bayes

02/01/2015 WEI PENG

a) Complete the function  $\text{logProd}(x)$  with takes as input a vector of numbers in logspace (i.e.,  $x_i = \log p_i$ ) and returns the product of those numbers in logspace-i.e.  $\text{logProd}(x) = \log(\prod_i p_i)$ .

```
function [log_product] = logProd(x)
% LOGPROD Given a vector of numbers in log-space
% (i.e. x(i) = log p(i)), this computes the log of the
% product of the numbers: log \product_i p(i).
%
% inputs:
%   x - A vector of length n containing the factors in
%       log-space: x(i) = log p(i).
%
% output:
%   log_product - A scalar containing the log of the
%       product of the numbers: log \product_i p(i).

% TODO: implement me!
log_product = 0;
log_product=sum(x);
end
```

b) Complete the function  $[D] = \text{NB\_XGivenY}(X\text{Train}, y\text{Train})$ . The output  $D$  is a  $2 \times V$  matrix, where for any word index  $\omega \in 1, \dots, V$  and class index  $y \in 0, 1$ , the entry  $D(y, \omega)$  is the MAP estimator of  $\theta_{y\omega} = P(X_\omega = 1 | Y = y)$  with a Beta(1,2) prior distribution.

```
function [D] = NB_XGivenY(XTrain, yTrain)
% NP_XGIVENY Estimates the probability that a word is
% observed given (conditioned on) the class label.
%
% inputs:
%   XTrain - An [n x V] matrix where each row describes
%       which words are present in a particular document.
%       XTrain(i, j) is 1 if word j appears in the i-th
%       document.
%   yTrain - An [n x 1] vector containing the class labels
%       for the training documents. yTrain(i) is 0 if the
%       i-th document belongs to The Economist or 1 if it
%       belongs to The Onion.
%
% output:
%   D - A [2 x V] matrix, where D(i, j) is an estimate of
%       the probability that word j is in a document with class
%       label i-1:  $P[X_j = 1 | Y = i-1]$ .

V = size(XTrain, 2);
D = zeros(2, V);
n=size(XTrain, 1);
N0=0;N1=0;
for i=1:n
    if yTrain(i)==0
        N0=N0+1;
        for j=1:V
            if XTrain(i, j)==1
                D(1, j)=D(1, j)+1;
            end
        end
    end
end
```

```

    else
        N1=N1+1;
        for j=1:V
            if XTrain(i,j)==1
                D(2,j)=D(2,j)+1;
            end
        end
    end
end
end
D(1,:)=(D(1,:)+1)/(N0+1);
D(2,:)=(D(2,:)+1)/(N1+1);
% TODO: implement me!
end

```

c) Complete the function  $[p] = NB_YPrior(yTrain)$ . The output  $p$  is the MLE for  $\rho = P(Y = 1)$ .

```

function [p] = NB_YPrior(yTrain)
%NP_YPRIOR Computes the prior probability that the class
% label is 1: P[Y = 1].
%
% inputs:
%   yTrain - An [n x 1] vector containing the class labels
%           for the training documents. yTrain(i) is 1 if the
%           i-th document belongs to The Economist or 2 if it
%           belongs to The Onion.
%
% output:
%   p - The prior probability that the class label is 1.
%
% TODO: implement me!
p = 0;
n=size(yTrain,1);
N0=0; N1=0;
for i=1:n
    if yTrain(i)==0
        N0=N0+1;
    else N1=N1+1;
    end
end
p(1)=N0/n;
p(2)=N1/n;
p=transpose(p);
end

```

d) Complete the function  $[yHat] = NB_cLASSIFY(D, P, X)$ . The input  $X$  is an  $n \times V$  matrix containing  $n$  feature vectors (stored as rows). The output  $yHat$  is an  $n \times 1$  vector of predicted class labels, where  $yHat(i)$  is the predicted label for the  $i^{th}$  row of  $X$ . (Hint: In this function, you will want to use the *logProd* function to avoid numerical problems).

```

function [yHat] = NB_Classify(D, p, XTest)
% NP_CLASSIFY Predicts the class labels for some given
% test input.
%
% inputs:
%   D - A [2 x V] matrix, where D(i, j) is an estimate of
%       the probability that word j in a document with class
%       label i-1 (the output of NB_XGivenY).
%   p - A scalar characterizing the prior probability that
%       the class label is 1 (the output of NB_YPrior).
%   XTest - An [m x V] matrix where each row describes
%           which words are present in a particular *test*

```

```

%      document. XTest(i, j) is 1 if word j appears in the
%      i-th document.
%
% output:
%      yHat - An [m x 1] vector of predicted class labels,
%      where yHat(i) is the predicted class label for the
%      i-th document (the i-th row of XTest).
m = size(XTest, 1);
V=size(XTest,2);
% TODO: implement me!
yHat = zeros(m, 1);
for i=1:m
    C0=zeros(V+1,1);
    C1=zeros(V+1,1);
    C0(V+1)=log(p(1));
    C1(V+1)=log(p(2));
    for j=1:V
        if XTest(i,j)==1
            C0(j)=log(D(1,j));
            C1(j)=log(D(2,j));
        else
            C0(j)=log(1-D(1,j));
            C1(j)=log(1-D(2,j));
        end
    end
    if logProd(C0)>logProd(C1)
        yHat(i)=0;
    else
        yHat(i)=1;
    end
end
end
end

```

e) Complete the function `[error] = ClassificationError(yHat,yTruth)`, which takes two vectors of equal length and returns the proportion of entries that they agree on.

```

function [error] = ClassificationError(yHat, yTruth)
% CLASSIFICATIONERROR Computes the classification error.
%
% inputs:
%      yHat - An [m x 1] vector of predicted class labels,
%      where yHat(i) is the predicted class label for the
%      i-th document (the i-th row of XTest).
%      yTruth - An [m x 1] vector of the true class labels.
%
% output:
%      error - A scalar containing the proportion of entries
%      that the given vectors disagree on.
% TODO: implement me!
error = 0;
m=size(yTruth,1);
for i=1:m
    if yHat(i)~= yTruth(i)
        error=error+1;
    end
end
error=error/m;
end

```

**Questions** f) Train your classifier on the data contained in ***XTrain***, and ***yTrain*** by running

```
D = NB_X_Given_Y(X_train, y_train)
P = NB_Y_Prior(y_train)
```

Use the learned classifier to predict the labels for the article feature vectors in ***XTrain*** and ***XTest*** by running

```
yHatTrain = NB_Classify(D, P, XTrain)
yHatTest = NB_Classify(D, P, XTest)
```

Use the function *ClassificationError* to measure and report the training and testing error by running

```
trainError = ClassificationError(yHatTrain, yTrain)
testError = ClassificationError(yHatTest, yTest)
```

(1) How do the train and test errors compare? Explain any significant difference.

A:

trainError=0, testError=0.0207.

g) Repeat the steps from part f), but this time use the smaller training set *XTrainSmall* and *yTrainSmall*. Explain any difference between the train and test error in the question and in part (f) [Hint: When we have less training data, does the prior have more or less impact on our classifier?]

(1) Explain any differences between the train errors in this question and in part (f).

(2) Explain any differences between the test errors in this question and in part (f).

A:

(1) trainErrorSmall=0.0966>0=trainError

(2) testErrorSmall=0.2759>0.0207=testError

h) Finally we will try to interpret the learned parameters. Train your classifier on the data contained in *XTrain* and *yTrain*. For each class label  $y \in 0, 1$ , list the six words that the model says are most likely to occur in a document from class  $y$ . Also for each class label  $y \in 0, 1$ , list the six words  $\omega$  that maximize the following quantity:

$$\frac{P(X_\omega = 1 | Y = y)}{P(X_\omega = 1 | Y \neq y)}$$

```
function [W] = Max8(Vocabulary, D)
W=cell(4,8);
A1=zeros(1,26048);
A2=zeros(1,26048);
[~,A1]=sort(D(1,:));
[~,A2]=sort(D(2,:));
[~,A3]=sort(D(1,:)./D(2,:));
[~,A4]=sort(D(2,:)./D(1,:));
for i=1:1:8
    W(1,i)=Vocabulary(A1(26040+i));
    W(2,i)=Vocabulary(A2(26040+i));
    W(3,i)=Vocabulary(A3(26040+i));
    W(4,i)=Vocabulary(A4(26040+i));
end
end
```

(1) Give one word in the top eight highest probabilities  $P(X_\omega = 1 | Y = 0)$  (i.e. words most likely to appear in Economist articles)

A: 'is' 'that' 'a' 'and' 'of' 'in' 'the' 'to'

(2) Give one word in the top eight highest probabilities  $P(X_\omega = 1 | Y = 1)$  (i.e. words most likely to appear in Onion articles).

A: 'a' 'and' 'the' 'to' 'of' 'said' 'in' 'for'

(3) Give one word in the top eight that maximize the quantity in Equation for  $y = 0$ .

A: 'parliamentari' 'neighbour' 'labour' '1990s' 'centr' 'favour' 'reckon' 'organis'

(4) Give one word in the top eight that maximize the quantity in Equation for  $y = 1$ .

A: 'favorit' 'tuesday' 'coach' 'realiz' 'percent' 'monday' '5enlarg' '4enlarg'

(5) Which list of words describes the two classes better?

A: The second list of words describes the two classes better.